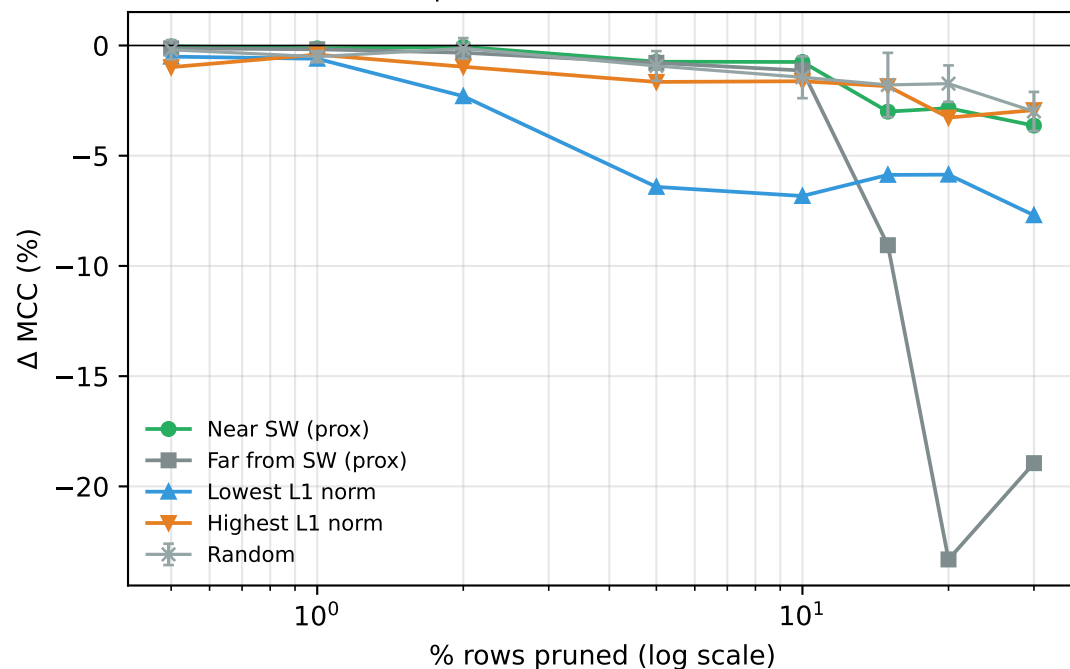
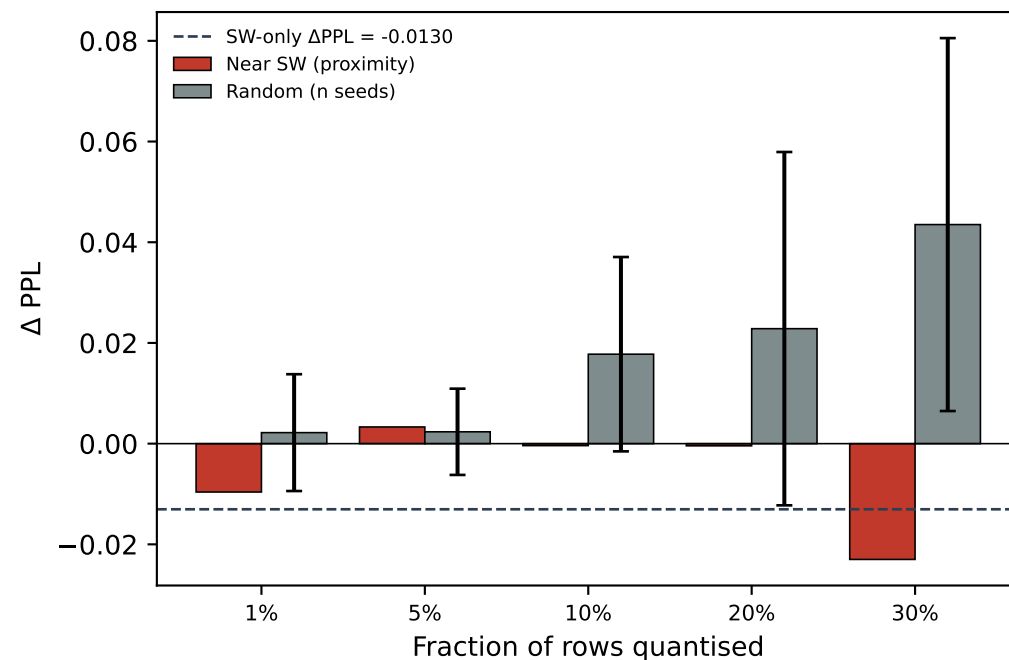


Figure 5 — Compression: shadow redundancy & whole-model INT4

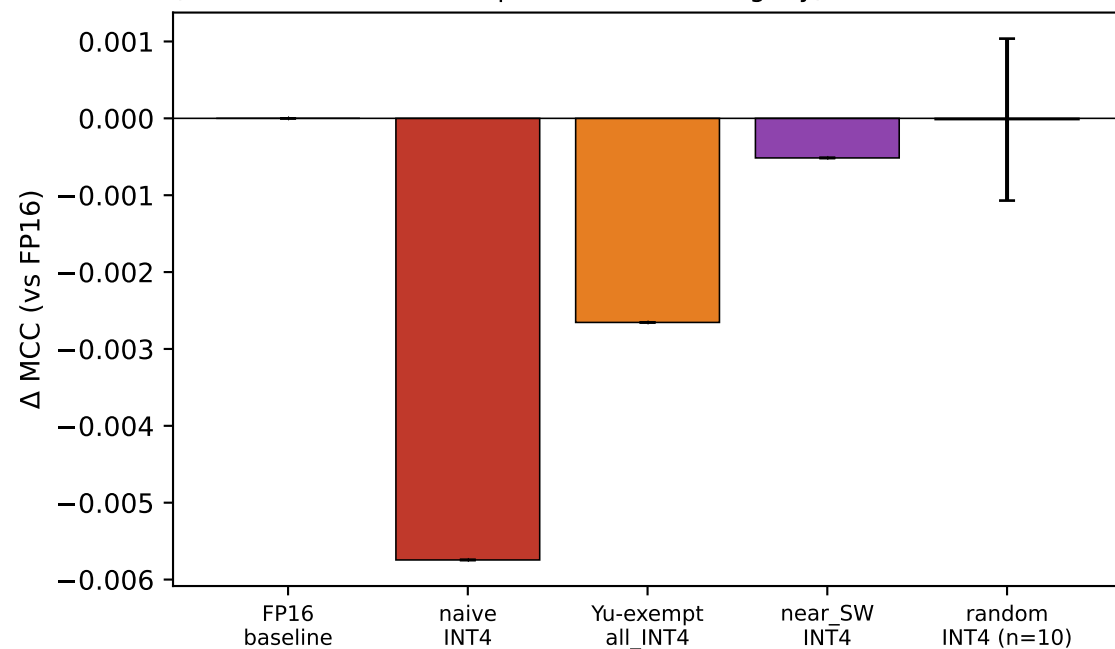
A. Pruning sweep — DNABERT-2 prom_core_notata
Shadow rows near SW prune at near-zero cost



B1. GENERator-EUK INT4 — per-row Δ PPL
(baseline PPL = 3.553)



B2. DNABERT-2 splice INT4 — downstream task
(near-SW $\approx$ random; no preferential SW fragility)



C. Whole-model INT4 — extended probe (100,000 tokens)

